

Coffee_Data_Analysis

April 3, 2023

1 Coffee Data Analysis and Visualization

1.1 By Dylan Benson

1.2 Setup

```
[2]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
#ML
from prophet import Prophet
sns.set(font_scale=1.5)
```

1.3 Read in and prepare data

```
[3]: pptg = pd.read_csv('./data/prices-paid-to-growers.csv')
rp = pd.read_csv('./data/retail-prices.csv')
cc = pd.read_csv('./data/Coffee-characteristics.csv')
cc = cc.drop(columns=['Farm.Name', 'Lot.Number', 'Certification.Address',
↪ 'Certification.Contact', 'Altitude', 'Region', 'Species', 'Mill', 'ICO.
↪ Number'])
dcc = pd.read_csv('./data/Domestic_Coffee_Consumption.csv')
dc = pd.read_csv('./data/domestic-consumption.csv')
tp = pd.read_csv('./data/total-production.csv')
```

```
[4]: #define directory paths

df_paths=[
    './data/domestic-consumption.csv',
    './data/exports-calendar-year.csv',
    './data/exports-crop-year.csv',
    './data/gross-opening-stocks.csv',
    './data/total-production.csv'
]

dfs=[pd.read_csv(df_path) for df_path in df_paths]
```

```

#define function making mean value of every column and attaching it to country
def get_means(df):
    df=df.copy()
    countries=df[df.columns[0]]
    mean=df.mean(axis=1)
    df=pd.concat([countries,mean],axis=1)
    df.columns=['country',countries.name]
    return df

```

```

[5]: #define function that creates data frames
def make_df(dfs):

    # Process all DataFrames
    processed_dfs = []
    for df in dfs:
        processed_dfs.append(get_means(df))

    # Merge DataFrames
    df = processed_dfs[0]

    for i in range(1, len(processed_dfs)):
        df = df.merge(processed_dfs[i], on='country')

    return df

data=make_df(dfs)

```

/tmp/ipykernel_85083/1234719787.py:17: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

```
mean=df.mean(axis=1)
```

/tmp/ipykernel_85083/1234719787.py:17: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

```
mean=df.mean(axis=1)
```

/tmp/ipykernel_85083/1234719787.py:17: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

```
mean=df.mean(axis=1)
```

/tmp/ipykernel_85083/1234719787.py:17: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

```
mean=df.mean(axis=1)
```

/tmp/ipykernel_85083/1234719787.py:17: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

```
mean=df.mean(axis=1)
```

```
[6]: ##rename columns and output to same csv (already done, doesn't need rerunning)
#df = df.rename(columns={'1990/91': '1990', '1991/92': '1991', '1992/93':
↳ '1992', '1993/94': '1993', '1994/95': '1994', '1995/96': '1995', '1996/97':
↳ '1996', '1997/98': '1997', '1998/99': '1998', '1999/00': '1999', '2000/01':
↳ '2000', '2001/02': '2001', '2002/03': '2002', '2003/04': '2003', '2004/05':
↳ '2004', '2005/06': '2005', '2006/07': '2006', '2007/08': '2007', '2008/09':
↳ '2008', '2009/10': '2009', '2010/11': '2010', '2011/12': '2011', '2012/13':
↳ '2012', '2013/14': '2013', '2014/15': '2014', '2015/16': '2015', '2016/17':
↳ '2016', '2017/18': '2017', '2018/19': '2018', '2019/20': '2019'})
#df.to_csv('Domestic_Coffee_Consumption.csv')

#Ensure no null values exist in our data
data.isna().sum() #returns False for all
data = data.dropna()

#Ensure no duplicate rows exist in our data
data.loc[data.duplicated()] #Nothing returned
data = data.drop_duplicates()

#reset data frame and index, sorting by domestic consumption
data = data.sort_values(by='domestic_consumption', ascending=False)
data = data.reset_index(drop=True)
data.head()
```

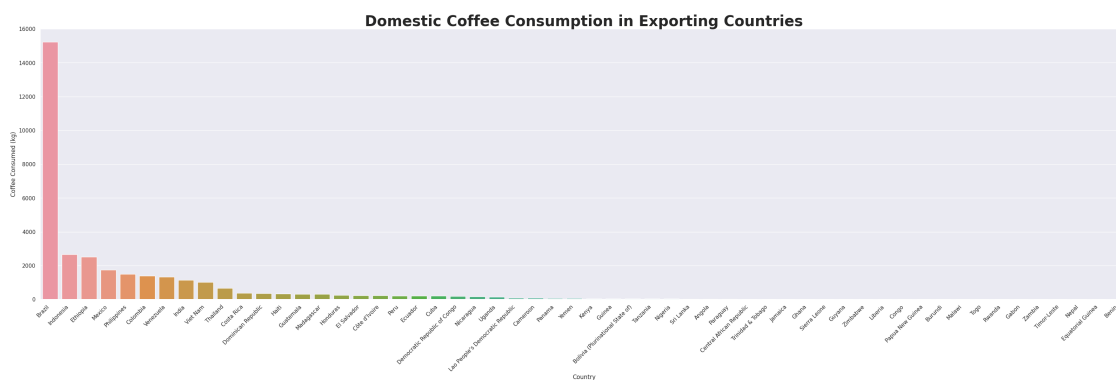
```
[6]:
```

	country	domestic_consumption	exports	exports_crop_year	\
0	Brazil	15234.310345	25706.195606	25919.128803	
1	Indonesia	2662.137931	5878.047357	5879.061059	
2	Ethiopia	2529.034483	2257.551574	2326.651490	
3	Mexico	1749.517241	3143.855086	3109.249886	
4	Philippines	1501.310345	29.944000	26.203414	


	gross_opening_stocks	total_production
0	23213.206897	41067.783976
1	690.114655	8452.302438
2	2044.586207	4880.789417
3	598.732759	4376.146438
4	563.620690	501.272379

1.4 Analysis and Visualization


```
[7]: sns.set(rc={"figure.figsize":(40, 10)})
consume_barplot = sns.barplot(x=data['country'], y =
    ↳data['domestic_consumption'])
consume_barplot.set_ylabel('Coffee Consumed (kg)')
consume_barplot.set_xlabel('Country')
consume_barplot.set_title('Domestic Coffee Consumption in Exporting Countries',
    ↳fontdict={'size': 30, 'weight': 'bold'})
consume_barplot.set_xticklabels(consume_barplot.get_xticklabels(), rotation=45,
    ↳horizontalalignment='right')
plt.ylim(0,16000)
plt.show()
```



```
[8]: #Create bar graph of top ten Countries by coffee consumption
top_ten_consume = data.head(10)
sns.set(rc={"figure.figsize":(40, 10)})
consume_topten_barplot = sns.barplot(x=top_ten_consume['country'], y =_
    ↳top_ten_consume['domestic_consumption'])
consume_topten_barplot.set_ylabel('Coffee Consumed (kg)')
consume_topten_barplot.set_xlabel('Country')
consume_topten_barplot.set_title('Domestic Coffee Consumption in Exporting_
    ↳Countries (Top Ten)', fontdict={'size': 30, 'weight': 'bold'})
plt.ylim(0,16000)
plt.show()
```




```
[9]: #Create bar graph of top ten Countries by coffee exports
data = data.sort_values(by='exports', ascending=False)
export_barplot = sns.barplot(x=data['country'], y = data['exports'])
export_barplot.set_ylabel('Coffee Exported (kg)')
export_barplot.set_xlabel('Country')
export_barplot.set_title('Coffee Exported by Country', fontdict={'size': 30,
    ↳'weight': 'bold'})
export_barplot.set_xticklabels(consume_barplot.get_xticklabels(), rotation=45,
    ↳horizontalalignment='right')
plt.ylim(0,27000)
plt.show()
```




```
[10]: #Create bar graph of top ten Countries by coffee exports
data = data.sort_values(by='exports', ascending=False)
top_ten_export = data.head(10)
export_topten_barplot = sns.barplot(x=top_ten_export['country'], y =
    ↳top_ten_export['exports'])
export_topten_barplot.set_ylabel('Coffee Exported (kg)')
export_topten_barplot.set_xlabel('Country')
export_topten_barplot.set_title('Coffee Exported by Country (Top Ten)',
    ↳fontdict={'size': 30, 'weight': 'bold'})
plt.ylim(0,30000)
```

```
plt.show()
```



```
[11]: #Drop the outlier in the data (Brazil)
no_brazil = data.drop(data.query("country=='Brazil'").index)
no_brazil = no_brazil.sort_values(by='domestic_consumption', ascending=False)
```

```
[12]: #Create same bar graphs, but excluding Brazil
top_ten_consume = no_brazil.head(10)
sns.set(rc={"figure.figsize":(40, 10)})
consume_barplot = sns.barplot(x=top_ten_consume['country'], y =
    ↳top_ten_consume['domestic_consumption'], dodge=False)
consume_barplot.set_ylabel('Coffee Consumed (kg)')
consume_barplot.set_xlabel('Year')
consume_barplot.set_title('Domestic Coffee Consumption in Exporting Countries_
    ↳(Excluding Brazil)', fontdict={'size': 30, 'weight': 'bold'})
plt.ylim(0,3000)
plt.show()
```




```
[13]: #Create bar graph of top ten Countries by coffee exports
no_brazil = no_brazil.sort_values(by='exports', ascending=False)
top_ten_export = no_brazil.head(10)
export_barplot = sns.barplot(x=top_ten_export['country'], y =
    ↳top_ten_export['exports'])
```

```

export_barplot.set_ylabel('Coffee Exported (kg)')
export_barplot.set_xlabel('Year')
export_barplot.set_title('Coffee Exported by Country (Excluding Brazil)',
    ↳fontdict={'size': 30, 'weight': 'bold'})
plt.ylim(0,14000)
plt.show()

```



1.5 Examine Correlation in the data

```

[14]: #Create heatmap of correlated data
stats = data[['exports', 'domestic_consumption', 'exports_crop_year',
    ↳'gross_opening_stocks', 'total_production']]
sns.set_theme(style="white")
corr = stats.corr(method = 'pearson', # The method of correlation
    ↳min_periods = 1 )
corr.style.background_gradient(cmap='coolwarm')


```

[14]: <pandas.io.formats.style.Styler at 0x7fa82b233b50>


```

[15]: #Scatter plot comparing domestic consumption x exports (excluding Brazil)
sns.set(rc={"figure.figsize":(20, 5)})
scatter = sns.scatterplot(data=no_brazil, x='domestic_consumption',
    ↳y='exports', legend='auto', s=50)
scatter.set_title("Comparing Domestic Coffee Consumption with Exports Among
    ↳Exporting Countries", fontdict={'size': 20, 'weight': 'bold'})
scatter.set_xlabel('Domestic Consumption (kg)', fontdict={'size': 15})
scatter.set_ylabel('Exports (kg)', fontdict={'size': 15})
plt.ylim(-1000, 14000)
plt.xlim(-100, 3000)
plt.show()

```



```
[16]: #Create a grid of pairplots between domestic consumption, exports, and
      ↪ production
      pairplot = sns.pairplot(data, vars=['domestic_consumption', 'exports',
      ↪ 'total_production'])
      plt.show()
```


```
[17]: #Create new dataframe of coffee types, clean data a bit (dcc =
      ↪Domestic_Coffee_Consumption.csv)
types = dcc['Coffee type']
types = types.replace({'Robusta/Arabica': 'Both'})
types = types.replace({'Arabica/Robusta': 'Both'})
pie = types.value_counts()
```

```
[18]: # Defining colors for the pie chart
colors = ['sienna', 'peru', 'cornsilk']

# Define the ratio of gap of each fragment in a tuple
x = (0.05, 0.05, 0.05)
```

```
#Create pie chart of coffee types
plt.ylabel(None)
pie.plot(kind='pie', title="Coffee Types Consumed in Exporting Countries",
        autopct='%1.0f%%', colors=colors, explode=x)
plt.ylabel(None)
plt.show()
```

Coffee Types Consumed in Exporting Countries



1.6 Examine Domestic Consumption Over Time

```
[19]: # (dc = domestic-consumption.csv)
#sort by consumption in 2018
dc = dc.sort_values(by='2018', ascending=False)
dc = dc.reset_index(drop=True)
top_ten = dc.head(10)

#Transpose the data frame
pivot = top_ten.transpose()

#rename columns to row 1
pivot.columns = pivot.iloc[0]
```

```

#drop first two rows
pivot = pivot.iloc[3:]

#rename index
pivot.index.names = ['Year']

copy = pivot.copy()


#Drop the outlier in the data (Brazil)
copy.drop('Brazil', axis=1, inplace=True)

top_ten_consume_overtime = copy.copy()

sns.set(rc={"figure.figsize":(30, 10)})

consume_plot = sns.lineplot(data=top_ten_consume_overtime, dashes=False)
consume_plot.set_title("Coffee Consumption Across Exporting Countries Over Time (Excluding Brazil)", fontdict={'size': 30, 'weight': 'bold'})
consume_plot.set_xlabel('Year', fontdict={'size': 15})
consume_plot.set_ylabel('Coffee Consumed (kg)', fontdict={'size': 15})
consume_plot.legend()
plt.ylim(0, 5000)
plt.show()

```



1.7 Examine Production over time

```

[20]: # (dc = domestic-consumption.csv)
#sort by consumption in 2018
tp = tp.sort_values(by='2018', ascending=False)
tp = tp.reset_index(drop=True)
top_ten_prod = tp.head(10)

```

```

#Transpose the data frame
pivot2 = top_ten_prod.transpose()

#rename columns to row 1
pivot2.columns = pivot2.iloc[0]

#drop first two rows
pivot2 = pivot2.iloc[3:]

#rename index
pivot2.index.names = ['Year']

copy = pivot2.copy()

#Drop the outlier in the data (Brazil)
copy.drop('Brazil', axis=1, inplace=True)

top_ten_produce_overtime = copy.copy()

sns.set(rc={"figure.figsize":(30, 10)})

produce_plot = sns.lineplot(data=top_ten_produce_overtime, dashes=False)
produce_plot.set_title("Coffee Production Across Exporting Countries Over Time_
↳(Excluding Brazil)", fontdict={'size': 30, 'weight': 'bold'})
produce_plot.set_xlabel('Year', fontdict={'size': 15})
produce_plot.set_ylabel('Coffee produced (kg)', fontdict={'size': 15})
produce_plot.legend()
plt.show()

```



1.8 Examine Percent change in consumption

```
[21]: dc_change = dc[dc['1990'] < dc['2018']]
dc_change['total_increase'] = dc_change['1990']/dc_change['2018']*100
inf = dc_change['total_increase'] == np.inf
dc_change.loc[inf, 'total_increase'] = 0
dc_change["total_increase"].fillna(0, inplace = True)
#dc_change['total_increase'].round(decimals = 2)
dc_change = np.round(dc_change, decimals = 2)
dc_change = dc_change.sort_values(by='total_increase', ascending=False)
increase_consume_barplot = sns.barplot(x=dc_change['domestic_consumption'], y =
    ↪dc_change['total_increase'])
increase_consume_barplot.set_ylabel('% Increase')
increase_consume_barplot.set_xlabel('Country')
increase_consume_barplot.set_title('Increase in Coffee Consumption among
    ↪Exporting Countries between 1990 and 2018', fontdict={'size': 30, 'weight':
    ↪'bold'})
increase_consume_barplot.set_xticklabels(increase_consume_barplot.
    ↪get_xticklabels(), rotation=45, horizontalalignment='right')
#plt.ylim(0,900)
plt.show()
```

/tmp/ipykernel_85083/3511321252.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy


```
dc_change['total_increase'] = dc_change['1990']/dc_change['2018']*100
```

/tmp/ipykernel_85083/3511321252.py:5: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
dc_change["total_increase"].fillna(0, inplace = True)
```



```
[22]: tp_change = tp[tp['1990'] < tp['2018']]
tp_change['total_increase'] = tp_change['1990']/tp_change['2018']*100
inf = tp_change['total_increase'] == np.inf
tp_change.loc[inf, 'total_increase'] = 0
tp_change["total_increase"].fillna(0, inplace = True)
tp_change = np.round(tp_change, decimals = 2)
tp_change = tp_change.sort_values(by='total_increase', ascending=False)
increase_produce_barplot = sns.barplot(x=tp_change['total_production'], y =
    ↳tp_change['total_increase'])
increase_produce_barplot.set_ylabel('% Increase')
increase_produce_barplot.set_xlabel('Country')
increase_produce_barplot.set_title('Increase in Coffee Production among
    ↳Exporting Countries between 1990 and 2018', fontdict={'size': 30, 'weight':
    ↳'bold'})
increase_produce_barplot.set_xticklabels(increase_produce_barplot.
    ↳get_xticklabels(), rotation=45, horizontalalignment='right')
plt.ylim(0,900)
plt.show()
```

/tmp/ipykernel_85083/1074522127.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

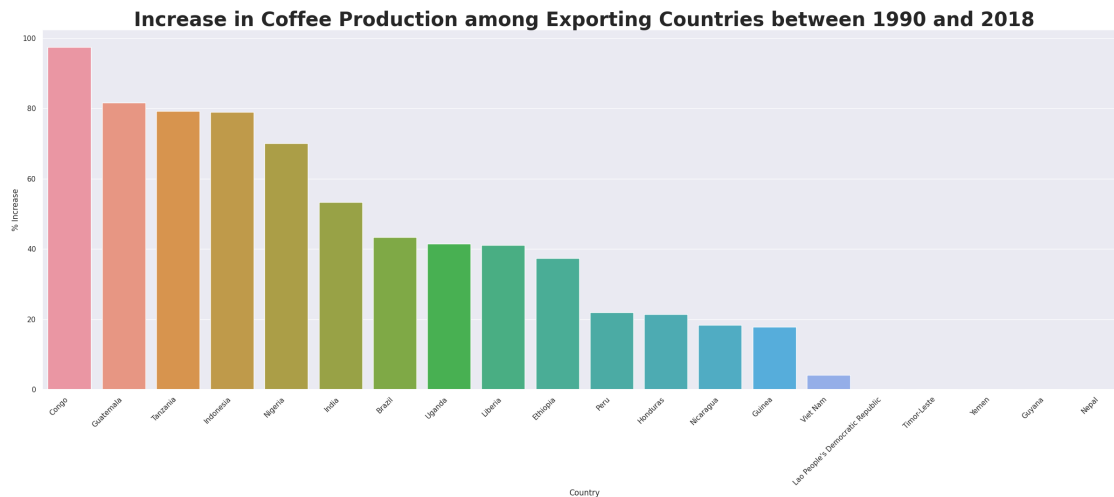
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
tp_change['total_increase'] = tp_change['1990']/tp_change['2018']*100
```

/tmp/ipykernel_85083/1074522127.py:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <https://pandas.pydata.org/pandas->

docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
 tp_change["total_increase"].fillna(0, inplace = True)



1.9 Analyze Brazil

```
[23]: #Reset pivot and create Brazil coffee consumption dataframe
pivot.drop(pivot.columns.difference(['Brazil']), 1, inplace=True)
Brazil_consume = pivot.copy()
Brazil_consume.rename(columns={'Brazil': 'consumption'}, inplace=True)

#Create another Brazil dataframe from the production data (tp =
↳ total-production.csv)
pivot2 = tp.transpose()
pivot2.columns = pivot2.iloc[0]
pivot2 = pivot2.drop('total_production')
pivot2.index.names = ['Year']
pivot2.drop(pivot2.columns.difference(['Brazil']), 1, inplace=True)
Brazil_prod = pivot2.copy()
Brazil_prod.rename(columns={'Brazil': 'production'}, inplace=True)
Brazil_prod.head()

#Combine the two
Brazil = pd.concat([Brazil_prod, Brazil_consume], axis=1)
Brazil.head(10)
```

/tmp/ipykernel_85083/1841956070.py:2: FutureWarning: In a future version of pandas all arguments of DataFrame.drop except for the argument 'labels' will be keyword-only.

```
    pivot.drop(pivot.columns.difference(['Brazil']), 1, inplace=True)
/tmp/ipykernel_85083/1841956070.py:11: FutureWarning: In a future version of
```


pandas all arguments of DataFrame.drop except for the argument 'labels' will be keyword-only.

```
pivot2.drop(pivot2.columns.difference(['Brazil']), 1, inplace=True)
```

```
[23]:      production consumption
```

```
Year
1990  27285.6286      NaN
1991  27293.4934      NaN
1992  34603.3542    8900.0
1993  28166.9786    9100.0
1994   28192.047    9300.0
1995  18060.2022   10100.0
1996   29196.743   11000.0
1997   26148.004   11500.0
1998   36760.8533  12200.0
1999   47577.8065  12700.0
```

```
[24]: #Create line graph of Brazilian coffee consumption and production over time.
sns.set(rc={"figure.figsize":(40, 10)})
brazil_consumption = sns.lineplot(data=Brazil, dashes=False)
brazil_consumption.set_title('Coffee Production x Consumption in Brazil Over_
    ↪Time', fontdict={'size': 30, 'weight': 'bold'})
brazil_consumption.set_xlabel('Year', fontdict={'size': 15})
brazil_consumption.set_ylabel('Coffee Consumed (billion kg)', fontdict={'size': 15,
    ↪15})
plt.ylim(0)
plt.show()
```



1.10 Examine Retail Prices vs Pay to Growers

```
[25]: #sort by retail price in 2018 (rp = retail-prices.csv)
rp = rp.sort_values(by='2018', ascending=False)
rp = rp.reset_index(drop=True)
rp.head()
```



```
[25]: retail_prices      1990      1991      1992      1993      1994 \
0 United Kingdom 23.289183 22.980132 22.273731 18.631347 25.077263
1      Italy 11.721854 12.406181 12.935982 10.132450 10.331126
2      Austria 10.816777 10.088300 11.015453 10.971302 10.110375
3      Japan 22.649007 26.225166 27.858720 32.163355 32.428256
4      Cyprus 6.247241 6.181015 6.335541 5.739514 7.019868

      1995      1996      1997      1998 ...      2009      2010 \
0 30.441501 29.470199 32.891832 34.039735 ... 35.298013 34.657837
1 12.582781 13.068433 12.030905 12.207506 ... 16.953642 16.203091
2 11.434879 11.964680 9.646799 8.763797 ... 15.342163 14.768212
3 39.116998 33.642384 31.390728 29.845475 ... 13.399558 14.105960
4 9.403974 9.116998 8.918322 10.176600 ... 12.207506 11.501104

      2011      2012      2013      2014      2015      2016 \
0 41.986755 42.384106 41.766004 45.386313 41.743929 35.960265
1 18.807947 18.741722 19.845475 19.536424 16.512141 16.445916
2 18.366446 18.498896 19.028698 19.050773 16.423841 12.450331
3 16.225166 16.710817 13.355408 12.538631 12.362031 13.708609
4 13.377483 14.039735 14.282561 14.304636 11.699779 11.699779

      2017      2018
0 37.549669 40.618102
1 16.931567 17.924945
2 13.730684 14.635762
3 13.134658 12.803532
4 12.141280 12.781457
```


[5 rows x 30 columns]

```
[26]: #Clean data and plot retail prices on line graph
pivot = rp.transpose()
pivot.columns = pivot.iloc[0]
pivot = pivot.iloc[1:]
pivot.index.names = ['Year']
pivot['Years'] = pivot.index
pivot['average'] = rp.mean()
retail_prices = sns.lineplot(data=pivot)
retail_prices.set_title('Retail Price of Coffee Across Importing Countries Over_
↳Time', fontdict={'size': 30, 'weight': 'bold'})
retail_prices.set_xlabel('Year', fontdict={'size': 15})
retail_prices.set_ylabel('Price of coffee per gram', fontdict={'size': 15})
retail_prices.legend()
plt.ylim(0, 50)
plt.show()
```

/tmp/ipykernel_85083/2299085409.py:7: FutureWarning: The default value of

numeric_only in DataFrame.mean is deprecated. In a future version, it will default to False. In addition, specifying 'numeric_only=None' is deprecated. Select only valid columns or specify the value of numeric_only to silence this warning.

```
pivot['average'] = rp.mean()
```



```
[27]: retail = rp[rp['1990'] < rp['2018']]
retail['total_increase'] = retail['1990']/retail['2018']*100
inf = retail['total_increase'] == np.inf
retail.loc[inf, 'total_increase'] = 0
retail["total_increase"].fillna(0, inplace = True)
retail = np.round(retail, decimals = 2)
retail = retail.sort_values(by='total_increase', ascending=False)
increase_rp_barplot = sns.barplot(x=retail['retail_prices'], y =
    ↪retail['total_increase'])
increase_rp_barplot.set_ylabel('% Increase')
increase_rp_barplot.set_xlabel('Country')
increase_rp_barplot.set_title('Change in Retail Price of Coffee among Importing
    ↪Countries between 1990 and 2018', fontdict={'size': 30, 'weight': 'bold'})
increase_rp_barplot.set_xticklabels(increase_rp_barplot.get_xticklabels(),
    ↪rotation=45, horizontalalignment='right')
plt.show()
```

/tmp/ipykernel_85083/2366707130.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead


See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
retail['total_increase'] = retail['1990']/retail['2018']*100
```

/tmp/ipykernel_85083/2366707130.py:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy


```
retail["total_increase"].fillna(0, inplace = True)
```




```
[28]: pivot2 = pptg.transpose()
pivot2.columns = pivot2.iloc[0]
pivot2 = pivot2.iloc[1:]
pivot2.index.names = ['Year']
pivot2['Years'] = pivot2.index
pivot2['average'] = pptg.mean()
grower_pay = sns.lineplot(data=pivot2)
grower_pay.set_title('Coffee Grower Pay Across Exporting Countries Over Time',
fontdict={'size': 30, 'weight': 'bold'})
grower_pay.set_xlabel('Year', fontdict={'size': 15})
grower_pay.set_ylabel('Pay per gram of Coffee', fontdict={'size': 15})
grower_pay.legend()
plt.ylim(0, 6)
plt.show()
```

/tmp/ipykernel_85083/461593401.py:6: FutureWarning: The default value of numeric_only in DataFrame.mean is deprecated. In a future version, it will default to False. In addition, specifying 'numeric_only=None' is deprecated. Select only valid columns or specify the value of numeric_only to silence this warning.

```
pivot2['average'] = pptg.mean()
```



```
[29]: pay_to_grow = pptg[pptg['1990'] < pptg['2018']]
pay_to_grow['total_increase'] = pay_to_grow['1990']/pay_to_grow['2018']*100
inf = pay_to_grow['total_increase'] == np.inf
pay_to_grow.loc[inf, 'total_increase'] = 0
pay_to_grow["total_increase"].fillna(0, inplace = True)
pay_to_grow = np.round(pay_to_grow, decimals = 2)
pay_to_grow = pay_to_grow.sort_values(by='total_increase', ascending=False)
increase_pptg_barplot = sns.barplot(x=pay_to_grow['prices_paid_to_growers'], y_
    ↪= pay_to_grow['total_increase'])
increase_pptg_barplot.set_ylabel('% Increase')
increase_pptg_barplot.set_xlabel('Country')
increase_pptg_barplot.set_title('Change in Pay To Coffee Growers among_
    ↪Exporting Countries between 1990 and 2018', fontdict={'size': 30, 'weight':_
    ↪'bold'})
increase_pptg_barplot.set_xticklabels(increase_pptg_barplot.get_xticklabels(),_
    ↪rotation=45, horizontalalignment='right')
plt.show()
```



```
[30]: #Combine average price to average pay over time
pivot.drop(pivot.columns.difference(['average']), 1, inplace=True)
avg_price = pivot.copy()
avg_price.rename(columns={'average': 'avg_price'}, inplace=True)

pivot2.drop(pivot2.columns.difference(['average']), 1, inplace=True)
avg_pay = pivot2.copy()
avg_pay.rename(columns={'average': 'avg_pay'}, inplace=True)

#Combine the two
compare_price_pay = pd.concat([avg_price, avg_pay], axis=1)
compare_price_pay = pd.DataFrame(compare_price_pay)
compare_price_pay.tail()
```


/tmp/ipykernel_85083/4150924749.py:2: FutureWarning: In a future version of pandas all arguments of DataFrame.drop except for the argument 'labels' will be

keyword-only.

```
pivot.drop(pivot.columns.difference(['average']), 1, inplace=True)
/tmp/ipykernel_85083/4150924749.py:6: FutureWarning: In a future version of
pandas all arguments of DataFrame.drop except for the argument 'labels' will be
keyword-only.
pivot2.drop(pivot2.columns.difference(['average']), 1, inplace=True)
```

```
[30]:      avg_price  avg_pay
Year
2014  15.439924  2.623876
2015  13.653422  2.255341
2016  12.658467  2.238045
2017  13.107852  2.279210
2018  13.651845  1.972795
```

```
[31]: #Plot the two
sns.set(font_scale=1.5)
price_pay_plot = sns.lineplot(data=compare_price_pay, dashes=False)
price_pay_plot.set_title('Average Price of Coffee x Average Pay to Grower Over_
↵Time', fontdict={'size': 30, 'weight': 'bold'})
price_pay_plot.set_xlabel('Year', fontdict={'size': 20})
price_pay_plot.set_ylabel('Price', fontdict={'size': 20})
plt.legend(loc='upper left', labels=['Price', 'Pay'])
plt.ylim(0,20)
plt.show()
```



2 Coffee Characteristics

```
[32]: # clean data (cc = Coffee-characteristics.csv)
cc = cc[pd.to_numeric(cc['ID'], errors='coerce').notnull()]
cc.head()
```

```
[32]:   ID      Owner Country.of.Origin \
0  1      metad plc      Ethiopia
1  2      metad plc      Ethiopia
```

2	3	grounds for health admin	Guatemala
3	4	yidnekachew dabessa	Ethiopia
4	5	metad plc	Ethiopia

	Company \
0	metad agricultural developmet plc
1	metad agricultural developmet plc
2	NaN
3	yidnekachew debessa coffee plantation
4	metad agricultural developmet plc

	Producer	Number.of.Bags	Bag.Weight \
0	METAD PLC	300	60 kg
1	METAD PLC	300	60 kg
2	NaN	5	1
3	Yidnekachew Dabessa Coffee Plantation	320	60 kg
4	METAD PLC	300	60 kg

	In.Country.Partner	Harvest.Year	Grading.Date ... \
0	METAD Agricultural Development plc	2014	April 4th, 2015 ...
1	METAD Agricultural Development plc	2014	April 4th, 2015 ...
2	Specialty Coffee Association	NaN	May 31st, 2010 ...
3	METAD Agricultural Development plc	2014	March 26th, 2015 ...
4	METAD Agricultural Development plc	2014	April 4th, 2015 ...

	Category.One.Defects	Quakers	Color	Category.Two.Defects	Expiration \
0	0	0	Green	0.0	April 3rd, 2016
1	0	0	Green	1.0	April 3rd, 2016
2	0	0	NaN	0.0	May 31st, 2011
3	0	0	Green	2.0	March 25th, 2016
4	0	0	Green	2.0	April 3rd, 2016

	Certification.Body	unit_of_measurement \
0	METAD Agricultural Development plc	m
1	METAD Agricultural Development plc	m
2	Specialty Coffee Association	m
3	METAD Agricultural Development plc	m
4	METAD Agricultural Development plc	m

	altitude_low_meters	altitude_high_meters	altitude_mean_meters
0	1950.0	2200.0	2075.0
1	1950.0	2200.0	2075.0
2	1600.0	1800.0	1700.0
3	1800.0	2200.0	2000.0
4	1950.0	2200.0	2075.0

[5 rows x 35 columns]

```

[33]: countries = cc.copy()
countries.drop(countries.columns.difference(['Country.of.Origin']), 1,
             inplace=True)
country_counts = countries.value_counts()
country_counts = pd.DataFrame(country_counts)
country_counts = country_counts.reset_index()
country_counts.columns=['Origin', 'Count']

#the top 5
country_counts2 = country_counts[:15].copy()

#others
new_row = pd.DataFrame(data = {
    'Origin' : ['Other'],
    'Count' : [country_counts['Count'][15:].sum()]
})

#combining top 5 with others
pie = pd.concat([country_counts2, new_row])

sns.set(rc={"figure.figsize":(40, 12)})

#define colors
colors = ['sienna', 'saddlebrown', 'chocolate', 'sandybrown', 'peru',
         'peachpuff', 'linen', 'seashell']

def autopct_format(values):
    def my_format(pct):
        total = sum(values)
        val = int(round(pct*total/100.0))
        return '{:.1f}%\n({v:d})'.format(pct, v=val)
    return my_format

pie_chart = pie.plot(kind = 'pie', y = 'Count', labels = pie['Origin'],
                    colors=colors, autopct=autopct_format(pie))
pie_chart.set_title('Coffee Country of Origin', fontdict={'size': 30, 'weight':
                    'bold'})
plt.legend([], [], frameon=False)
plt.ylabel(None)
plt.show()

```


/tmp/ipykernel_85083/2219077466.py:2: FutureWarning: In a future version of pandas all arguments of DataFrame.drop except for the argument 'labels' will be keyword-only.

```

countries.drop(countries.columns.difference(['Country.of.Origin']), 1,
inplace=True)

```

Coffee Country of Origin



2.1 Predict average production over next few years

2.2 Pay

```
[34]: predict_pay = avg_pay.rename(columns={'avg_pay': 'ds'})
      predict_pay['y'] = predict_pay.index
      predict_pay.tail()
```

```
[34]: prices_paid_to_growers      ds      y
      Year
      2014      2.623876  2014
      2015      2.255341  2015
```


2016	2.238045	2016
2017	2.279210	2017
2018	1.972795	2018

```
[35]: split_date = '2018'
pay_train = predict_pay.loc[predict_pay['y'] <= split_date].copy()
pay_test = predict_pay.loc[predict_pay['y'] > split_date].copy()

# Plot train and test so you can see where we have split
pay_test = pay_test.rename(columns={'ds': 'TEST SET'})
pay_train = pay_train.rename(columns={'ds': 'TRAINING SET'})
pay_set = pay_train.merge(pay_test, how = 'outer')
pay_set.index = pay_set['y']
#pay_set.plot(figsize=(10, 5), title='Avg Pay over time', style='.', ms=1)
pay_set_plot = sns.scatterplot(data=pay_set, s=150)
pay_set_plot.set_title('Coffee Grower Pay Across Exporting Countries Over_
↪Time', fontdict={'size': 30, 'weight': 'bold'})
plt.ylim(0, 5)
pay_set_plot.set_xlabel('Year', fontdict={'size': 15})
pay_set_plot.set_ylabel('Pay per gram of Coffee', fontdict={'size': 15})
pay_set_plot.legend()
#plt.xlim(1990, 2018)
plt.show()
```



```
[36]: # Format data for prophet model using ds and y
pay_train_prophet = pay_train.reset_index() \
    .rename(columns={'y': 'ds',
                    'TRAINING SET': 'y'})

pay_train_prophet.head()
```

```
[36]: prices_paid_to_growers  Year      y      ds
0                1990  1.026493  1990
1                1991  1.018779  1991
```

2	1992	0.847206	1992
3	1993	0.898745	1993
4	1994	1.769350	1994

```
[37]: %%time
model = Prophet()
model.fit(pay_train_prophet)
```

16:47:34 - cmdstanpy - INFO - Chain [1] start processing

16:47:34 - cmdstanpy - INFO - Chain [1] done processing

CPU times: user 167 ms, sys: 16.7 ms, total: 184 ms


Wall time: 368 ms

```
[37]: <prophet.forecaster.Prophet at 0x7fa82a4f3fd0>
```

```
[38]: #Predict the future
future = model.make_future_dataframe(periods=20, freq='y',
    ↪include_history=False)
forecast = model.predict(future)
```

```
[39]: fig, ax = plt.subplots(figsize=(10, 5))
ax.scatter(forecast['ds'], forecast['yhat'], color='r')
fig = model.plot(forecast, ax=ax)
#ax.set_ylim(0, 70000)
plot = plt.suptitle('Predicting Future Pay to Coffee Growers')
ax.set_xlabel("Year")
ax.set_ylabel("Pay per gram of Coffee")
```

```
[39]: Text(82.0, 0.5, 'Pay per gram of Coffee')
```




2.3 Price

```
[40]: predict_price = avg_price.rename(columns={'avg_price': 'ds'})
      predict_price['y'] = predict_price.index
      predict_price.tail()
```

```
[40]: retail_prices      ds      y
      Year
      2014      15.439924  2014
      2015      13.653422  2015
      2016      12.658467  2016
      2017      13.107852  2017
      2018      13.651845  2018
```

```
[41]: split_date = '2018'
      price_train = predict_price.loc[predict_price['y'] <= split_date].copy()
      price_test = predict_price.loc[predict_price['y'] > split_date].copy()

      # Plot train and test so you can see where we have split
      price_test = price_test.rename(columns={'ds': 'TEST SET'})
      price_train = price_train.rename(columns={'ds': 'TRAINING SET'})
      price_set = price_train.merge(price_test, how = 'outer')
      price_set.index = price_set['y']
      price_set_plot = sns.scatterplot(data=price_set, s=150)
      price_set_plot.set_title('Coffee Grower price Across Exporting Countries Over_
      ↪Time', fontdict={'size': 30, 'weight': 'bold'})
      plt.ylim(0, 20)
      price_set_plot.set_xlabel('Year', fontdict={'size': 15})
      price_set_plot.set_ylabel('price per gram of Coffee', fontdict={'size': 15})
      price_set_plot.legend()
      #plt.xlim(1990, 2018)
      plt.show()
```



```
[42]: # Format data for prophet model using ds and y
price_train_prophet = price_train.reset_index() \
    .rename(columns={'y':'ds',
                    'TRAINING SET':'y'})

price_train_prophet.head()
```

```
[42]: retail_prices  Year      y    ds
0          1990  10.386313  1990
1          1991  10.424156  1991
2          1992  10.532955  1992
3          1993   9.651529  1993
4          1994  10.994954  1994
```

```
[43]: %%time
model2 = Prophet()
model2.fit(price_train_prophet)
```

```
16:47:36 - cmdstanpy - INFO - Chain [1] start processing
16:47:36 - cmdstanpy - INFO - Chain [1] done processing

CPU times: user 32.7 ms, sys: 11.2 ms, total: 43.9 ms
Wall time: 178 ms
```

```
[43]: <prophet.forecaster.Prophet at 0x7fa82831ee30>
```

```
[44]: #Predict the future
future2 = model2.make_future_dataframe(periods=20, freq='y',
    ↪include_history=False)
forecast2 = model2.predict(future2)
```

```
[45]: fig, ax = plt.subplots(figsize=(10, 5))
ax.scatter(forecast2['ds'], forecast2['yhat'], color='r')
fig = model2.plot(forecast2, ax=ax)
#ax.set_ylim(0, 70000)
plot = plt.suptitle('Predicting Future Price of Coffee')
ax.set_xlabel("Year")
ax.set_ylabel("Price per gram of Coffee")
plt.show()
```

